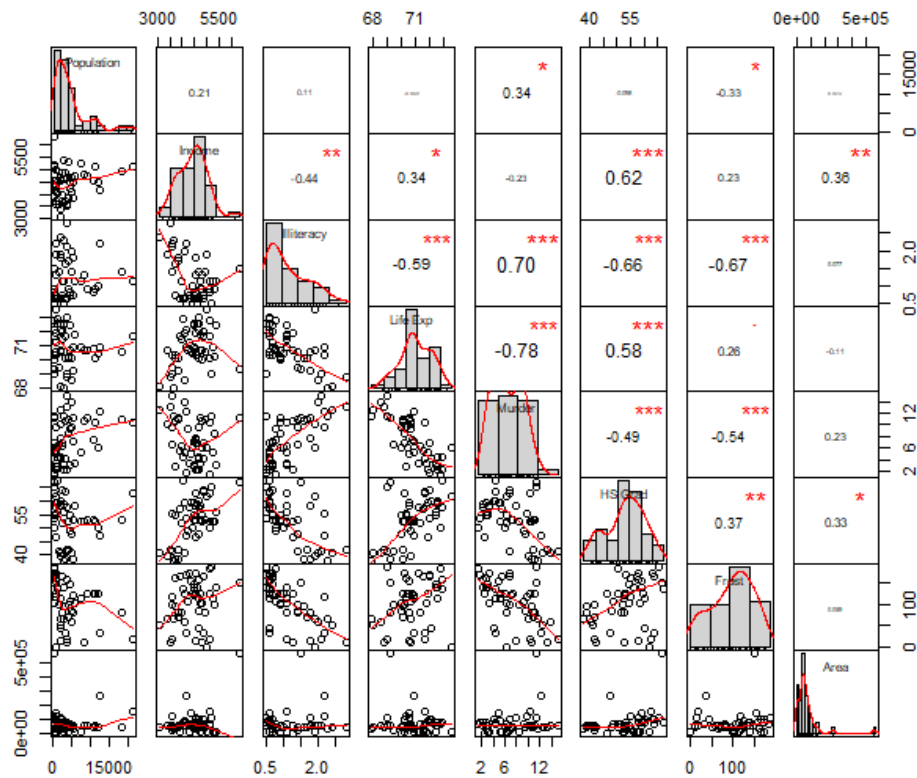


Problem 1

- a) `state.x77` dataset has 50 observations on 8 variables. There is no respond variable and all variables are considered as predictor variables. All the variables are quantitative. Bellow is a summary of these variables.

Population	Income	Illiteracy	Life Exp	Murder	HS Grad
Min. : 365	Min. :3098	Min. :0.500	Min. :67.96	Min. : 1.400	Min. :37.80
1st Qu.: 1080	1st Qu.:3993	1st Qu.:0.625	1st Qu.:70.12	1st Qu.: 4.350	1st Qu.:48.05
Median : 2838	Median :4519	Median :0.950	Median :70.67	Median : 6.850	Median :53.25
Mean : 4246	Mean :4436	Mean :1.170	Mean :70.88	Mean : 7.378	Mean :53.11
3rd Qu.: 4968	3rd Qu.:4814	3rd Qu.:1.575	3rd Qu.:71.89	3rd Qu.:10.675	3rd Qu.:59.15
Max. :21198	Max. :6315	Max. :2.800	Max. :73.60	Max. :15.100	Max. :67.30
Frost	Area				
Min. : 0.00	Min. : 1049				
1st Qu.: 66.25	1st Qu.: 36985				
Median :114.50	Median : 54277				
Mean :104.46	Mean : 70736				
3rd Qu.:139.75	3rd Qu.: 81163				
Max. :188.00	Max. :566432				

Following figure shows how most of these variables, specially Life Exp, Murder, HS Grad and Frost, are highly correlates with other variables.





- b) Data needs to be standardized in order to make the sample mean of principle components = 0 and variance equal to the eigen values of them. Thus make them uncorrelated.

```
> apply(state.x77, 2, mean)
Population      Income Illiteracy    Life Exp      Murder      HS Grad      Frost      Area
4246.4200  4435.8000    1.1700    70.8786    7.3780    53.1080   104.4600  70735.8800

> apply(state.x77, 2, sd)
Population      Income Illiteracy    Life Exp      Murder      HS Grad      Frost
4.464491e+03  6.144699e+02  6.095331e-01  1.342394e+00  3.691540e+00  8.076998e+00  5.198085e+01
Area
8.532730e+04
```

- c) PCA was preformed on the data and proportion of variance explained (PVE) was computed to find the number of PCs needed. Bellow are the graphs of PVE and cumulative of PVE against the number of PCs.



```
> cumsum(pve)
```

```
[1] 0.4498619 0.6538519 0.7928445 0.8812825 0.9293627 0.9677954 0.9858515 1.0000000
```

In the first graph there is a big drop from 1PC to 2PCs and a small drop from 6PCs to 7PCs. Also by looking at the Cumulative of PVE, we have 65% with 2PCs and 92% with 5PCs. So I would recommend to have more than 5PCs for this data set.


d)

	Population	Income	Illiteracy	Life Exp	Murder	HS Grad	Frost	Area
PC1	0.2398436	-0.5669029	0.88720374	-0.7808560	0.8427885	-0.8056584	-0.6780384	-0.06333314
PC2	0.5248778	0.6629778	0.06766573	-0.1043129	0.3921173	0.3816642	-0.1961984	0.75067024

cumulative.PVE

PC1	0.4498619
PC2	0.6538519

With the figure, its clear that the two components are related to most of the states. This can also be seen with the cumulative percentage of the total variability explained. When you have the first two PCs cumulative.PVE = 0.6538519 and thus they were enough to explain more than half of the data. Most of the unexplained states are southern. Thus there is a southern component in the figure.



Problem 2

a)

```
> signif(sort(train.latent.sem[, 1], decreasing = TRUE)[1:30], 2)
theater      music      m      theaters    composers    matinee      opera
0.180      0.130      0.080      0.079      0.078      0.077      0.075
> signif(sort(train.latent.sem[, 1], decreasing = FALSE)[1:30], 2)
her      she      ms      painting    paintings    mother      cooper      artists
-0.150    -0.140    -0.130    -0.110    -0.100    -0.092    -0.090    -0.086
> signif(sort(train.latent.sem[, 2], decreasing = TRUE)[1:30], 2)
she      her      theater    said      i      ms      mother      cooper
0.240    0.240    0.200    0.170    0.120    0.110    0.110    0.110
> signif(sort(train.latent.sem[, 2], decreasing = FALSE)[1:30], 2)
patterns    chinese    feb      chelsea    computers    diamond    europeans
-0.065    -0.051    -0.046    -0.046    -0.046    -0.045    -0.044
```

With the 1st PC, words with positive projections are mostly associated with music, those with negative components with the visual arts and with the 2nd PC, the positive words are about art and the negative words are musical.





Figure : Projection of the Times stories on to the first two principal components. Music stories are marked with a blue “m”, art stories with a red “a”.

Total number of PCs is 80 and this is what we would expect since the training data set consist of 80 observations and 4432 predictors.

- b) By looking at the above figure, we can see that almost all the words are separated into two groups, arts and music, with few been scattered around. So we can use the first two PCs for a good prediction.

c) estimate of the training error rate = 0.1625



With the glm fit, we were able to give a decision boundary for the separation of the words into arts and music. Above graph shows this linear boundary and most of the words can be identify as arts or music with an error rate of 0.1625.

d)

```
test.predicted.classes
  art music
art   10   2
music  5   5
```

test error rate = $(5+2) / 22 = 0.3181818$

class specific error rates

art = $2/12 = 0.1666667$

music = $5/10 = 0.5$






Figure shows how the 22 words in the test data set is scattered around and how the decision boundary from the glm fit is able to classify them.

e) We can increase the number of PCs considered for the glm fit and check the test error rate with the test data set so than we can pick M as the number of PCs that give the minimum test error rate. PCR method is used when the respond variable is qualitative. Thus its best to use PCA instead of PCR.

Problem 3



a) Test error = 1276986

b) Optimal $\lambda = 0.01$



Test error = 1277006

c) Optimal $\lambda = 0.01$




Test error = 1277099

d)

> MSEP(pcr.fit)

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps
CV	1.5e+07	14744421	4123358	4157777	2883584	2506292	2510774	2492036	2376917	2245553
adjCV	1.5e+07	14744466	4123261	4157770	2879623	2505933	2510639	2492027	2376694	2245437
	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	16 comps	17 comps		
CV	2237022	2251874	2251604	2267449	2267624	2067302	1357315	1276987		
adjCV	2236914	2251762	2251488	2267334	2267512	2066763	1357131	1276825		

17PCs gives them minimum MSE. So the test error = 1276987




e)

```
> MSEPls(pls.fit)
```

	(Intercept)	1 comps	2 comps	3 comps	4 comps	5 comps	6 comps	7 comps	8 comps	9 comps
CV	1.5e+07	3418008	2414639	2049701	1806343	1371151	1325548	1301942	1297801	1283053
adjCV	1.5e+07	3417887	2414598	2049563	1805870	1370529	1325245	1301788	1297869	1282936

	10 comps	11 comps	12 comps	13 comps	14 comps	15 comps	16 comps	17 comps
CV	1282432	1280546	1279732	1278057	1277875	1277232	1276966	1276987
adjCV	1282273	1280398	1279576	1277894	1277712	1277070	1276804	1276825

16PCs gives them minimum MSEP. So the test error = 1276966



	Linear model	Ridge	Lasso	PCR	PLS
Test error (MSEP)	1276986	1277006	1277099	1276987	1276966

Although all the models gives approximately closer test errors, lowest test error gives from PLS model.

```

# Problem 1

head(state.x77)
str(state.x77)
# Extract the names of states
states <- row.names(state.x77)

# part a)
summary(state.x77)
chart.Correlation(state.x77) #matrix of scatterplot and correlation after transformation

# part b)
# Look at mean and sd
apply(state.x77, 2, mean)
apply(state.x77, 2, sd)

# part c)
# Perform PCA
pca <- prcomp(state.x77, center = T, scale = T)
names(pca)

# Get the loading matrix
pca$rotation

# Get the score matrix
dim(pca$x)
head(pca$x)

# Check the covariance matrix of the scores
round(cov(pca$x), 4)

# Display a biplot the results (shows both pc scores and loading vectors)
pca$rotation
biplot(pca, scale=0)

# Display the biplot after changing the signs of loadings and scores
pca$rotation <- -pca$rotation
pca$x <- -pca$x
pca$rotation
biplot(pca, scale=0)

# Compute the proportion of variance explained (PVE)
pc.var <- pca$sdev^2
pve <- pc.var/sum(pc.var)
pve
cumsum(pve)

#Correlations
rbind(PC1 = c(pca$rotation[,1]*pca$sdev[1], cumulative.PVE = cumsum(pve)[1]),
      PC2 = c(pca$rotation[,2]*pca$sdev[2], cumulative.PVE = cumsum(pve)[2]))

par(mfrow=c(1,2))
# Scree plot
plot(pve, xlab = "Principal Component", ylab = "Proportion of Variance Explained", ylim =
c(0,1), type = 'b')

# Plot of cumulative PVE
plot(cumsum(pve), xlab = "Principal Component", ylab = "Cumulative Proportion of Variance
Explained", ylim = c(0,1), type = 'b')
par(mfrow=c(1,1))

```

```

#Correlations
rbind(PC1 = c(pca$rotation[,1]*pca$sdev[1], cumulative.PVE = cumsum(pve)[1]),
      PC2 = c(pca$rotation[,2]*pca$sdev[2], cumulative.PVE = cumsum(pve)[2]))

# problem 2

train.data <- read.csv("nyt.train.csv", header = T)
test.data <- read.csv("nyt.test.csv", header = T)

train.pca <- prcomp(train.data[, -1])
# We need to omit the first column because it contains categorical variables, and PCA
doesn't apply to them.
train.latent.sem <- train.pca$rotation
signif(sort(train.latent.sem[, 1], decreasing = TRUE)[1:30], 2)
# theater music m theaters composers matinee opera
# 0.180 0.130 0.080 0.079 0.078 0.077 0.075
# sunday musical jersey p orchestra band committee
# 0.067 0.065 0.064 0.064 0.062 0.061 0.060
# performance performances east organ dance hour program
# 0.059 0.056 0.056 0.053 0.052 0.051 0.051
# events yesterday will recitals ballet purchase X.d
# 0.050 0.049 0.049 0.048 0.048 0.048 0.047
# guitarist calif
# 0.045 0.044

signif(sort(train.latent.sem[, 1], decreasing = FALSE)[1:30], 2)
# her she ms painting paintings mother cooper artists
# -0.150 -0.140 -0.130 -0.110 -0.100 -0.092 -0.090 -0.086
# white images i said process sculpture picasso gagosian
# -0.078 -0.077 -0.071 -0.070 -0.070 -0.070 -0.068 -0.065
# art my nature image color sculptures work red
# -0.064 -0.064 -0.064 -0.061 -0.061 -0.059 -0.059 -0.058
# artist rothko paint photographs paper figure
# -0.056 -0.055 -0.055 -0.055 -0.054 -0.054

signif(sort(train.latent.sem[, 2], decreasing = TRUE)[1:30], 2)
# she her theater said i ms mother cooper
# 0.240 0.240 0.200 0.170 0.120 0.110 0.110 0.110
# says opera my hour id im production was
# 0.089 0.084 0.084 0.082 0.081 0.079 0.075 0.075
# mrs play sir broadway awards you national garde
# 0.074 0.074 0.071 0.070 0.066 0.066 0.065 0.063
# me season jonathan week baby networks
# 0.062 0.062 0.062 0.060 0.059 0.059

signif(sort(train.latent.sem[, 2], decreasing = FALSE)[1:30], 2)
# patterns chinese feb chelsea computers diamond europeans
# -0.065 -0.051 -0.046 -0.046 -0.046 -0.045 -0.044
# gallery museum art heads white stone views
# -0.042 -0.041 -0.040 -0.039 -0.039 -0.039 -0.039
# painted recalling soho artists pills statue newman
# -0.039 -0.039 -0.039 -0.038 -0.037 -0.037 -0.037
# computer compositions grid landscapes spatial images wood
# -0.037 -0.037 -0.037 -0.037 -0.037 -0.036 -0.035
# technology personal
# -0.035 -0.035

```



```

plot(train.pca$x[, 1:2],
     pch = ifelse(train.data[, "class.labels"] == "music", "m", "a"),
     col = ifelse(train.data[, "class.labels"] == "music", "blue", "red"))

# part c)
train.set <- data.frame(class.labels=train.data$class.labels,
                        CP1 = train.pca$x[, 1],
                        CP2 = train.pca$x[, 2])
predictors <- train.pca$x[, 1:2]
lm.fit <- glm(class.labels ~ CP1 + CP2, data = train.set, family = "binomial")

lm.pred <- predict(lm.fit, train.set, type = 'response')
predicted.classes <- as.factor(ifelse(lm.pred < 0.5, 'art', 'music'))
mean(train.data$class.labels != predicted.classes) # training error
# [1] 0.1625

slope <- coef(lm.fit)[2]/(-coef(lm.fit)[3])
intercept <- coef(lm.fit)[1]/(-coef(lm.fit)[3])

abline(intercept , slope)

# part d)
pca.scores <- predict(train.pca, test.data)[,1:2]
test.set <- cbind.data.frame(class.labels = test.data$class.labels,
                             CP1 = pca.scores[, 1],
                             CP2 = pca.scores[, 2])

plot(pca.scores[, 1:2],
     pch = ifelse(test.data[, "class.labels"] == "music", "m", "a"),
     col = ifelse(test.data[, "class.labels"] == "music", "blue", "red"))

test.lm.pred <- predict(lm.fit, test.set , type = 'response')
test.predicted.classes <- as.factor(ifelse(test.lm.pred < 0.5, 'art', 'music'))
mean(test.data$class.labels != test.predicted.classes) # test error
# [1] 0.3181818
abline(intercept , slope)

# confusion matrix for test and training data
table(test.data$class.labels, test.predicted.classes)

# problem 3

library(caret) # for cross-validation
library(ISLR)
library(glmnet) # for ridge and lasso
library(pls) # for pcr and pls
train.data <- College
str(train.data)
attach(train.data)

# part a)
lm.fit <- train(Apps ~ .,

```

```

        method = "lm",
        data = train.data,
        trControl = trainControl(method = "LOOCV"))

print(lm.fit)
# Resampling results:
#
#      RMSE      Rsquared  MAE
# 1130.038  0.91468   630.0335

1130.038^2 # = 1276986

lm.pred <- predict(lm.fit, train)
mean((Apps - lm.pred)^2) # training error
# [1] 1059279

# part b)
# Create response vector and the design matrix (without the first column of 1s)
y <- Apps
x <- model.matrix(Apps ~ ., train.data)[, -1]

grid <- 10^seq(8, -2, length = 200)

# Fit ridge regression for each lambda on the grid
ridge.out <- glmnet(x, y, alpha = 0, lambda = grid)
plot(ridge.out, xvar = "lambda")

# leave one out cross-validation
set.seed(1)
ridge.cv.out <- cv.glmnet(x, y, alpha = 0, lambda = grid, nfolds = dim(train.data)[1])
plot(ridge.cv.out)

print(ridge.cv.out)
# Measure: Mean-Squared Error
#
#      Lambda Measure      SE Nonzero
# min    0.01 1277006 258548      17
# 1se 235.43 1516697 508117      17

# Find the best value of lambda
ridge.bestlam <- ridge.cv.out$lambda.min
# [1] 0.01
# log(ridge.bestlam) = -4.60517

coef.ridge <- predict(ridge.out, type = "coefficients", s = ridge.bestlam)
#
#      1
# (Intercept) -445.26830402
# PrivateYes -494.15980608
# Accept      1.58570739
# Enroll     -0.88022903
# Top10perc   49.92174051
# Top25perc  -14.23153321
# F.Undergrad  0.05734870
# P.Undergrad  0.04444658
# Outstate   -0.08586349
# Room.Board  0.15104313
# Books       0.02090569
# Personal    0.03109799

```

```
# PhD -8.67805769
# Terminal -3.33091855
# S.F.Ratio 15.38992988
# perc.alumni 0.17691111
# Expend 0.07789878
# Grad.Rate 8.66799704
```

```
ridge.pred <- predict(ridge.out, s = ridge.bestlam, newx = x, type='response')
train.error.ridge <- mean((ridge.pred - y)^2) # training error
# [1] 1059279
```

```
# part c)
grid <- 10^seq(5, -2, length = 200)
```

```
# Fit lasso regression for each lambda on the grid
lasso.out <- glmnet(x, y, alpha = 1, lambda = grid)
plot(lasso.out, xvar = "lambda")
```

```
# leave one out cross-validation
set.seed(1)
lasso.cv.out <- cv.glmnet(x, y, alpha = 1, lambda = grid, nfolds = dim(train.data)[1])
plot(lasso.cv.out)
```

```
print(lasso.cv.out)
# Measure: Mean-Squared Error
#
#      Lambda Measure      SE Nonzero
# min    0.01 1277099 258652      17
# 1se  318.1 1524047 355598       3
```

```
# Find the best value of lambda
lasso.bestlam <- lasso.cv.out$lambda.min
# [1] 0.01
# log(lasso.bestlam) = -4.60517
```

```
coef.lasso <- predict(lasso.out, type = "coefficients", s = lasso.bestlam)
#
# (Intercept) -445.26830402
# PrivateYes -494.15980608
# Accept 1.58570739
# Enroll -0.88022903
# Top10perc 49.92174051
# Top25perc -14.23153321
# F.Undergrad 0.05734870
# P.Undergrad 0.04444658
# Outstate -0.08586349
# Room.Board 0.15104313
# Books 0.02090569
# Personal 0.03109799
# PhD -8.67805769
# Terminal -3.33091855
# S.F.Ratio 15.38992988
# perc.alumni 0.17691111
# Expend 0.07789878
# Grad.Rate 8.66799704
```

```

lasso.pred <- predict(lasso.out, s = lasso.bestlam, newx = x, type='response')
train.error.lasso <- mean((lasso.pred - y)^2) # training error
# [1] 1059279

# part d)
pcr.fit <- pcr(Apps~., data=train.data, scale=T, validation="LOO")

# Scree plot
validationplot(pcr.fit, val.type="MSEP")

MSEP(pcr.fit)
#      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8
comps  9 comps 10 comps
# CV      1.5e+07 14744421 4123358 4157777 2883584 2506292 2510774 2492036
2376917 2245553 2237022
# adjCV    1.5e+07 14744466 4123261 4157770 2879623 2505933 2510639 2492027
2376694 2245437 2236914
#      11 comps 12 comps 13 comps 14 comps 15 comps 16 comps 17 comps
# CV      2251874 2251604 2267449 2267624 2067302 1357315 1276987
# adjCV    2251762 2251488 2267334 2267512 2066763 1357131 1276825

pcr.pred <- predict(pcr.fit, train.data, ncomp=which.min(MSEP(pcr.fit)$val[1,1,]) - 1)
mean((Apps - pcr.pred)^2) # training error
# [1] 1059279

# part e)
pls.fit <- pls(Apps~., data=train.data, scale=T, validation="LOO")

# Scree plot
validationplot(pls.fit, val.type="MSEP")

MSEP(pls.fit)
#      (Intercept)  1 comps  2 comps  3 comps  4 comps  5 comps  6 comps  7 comps  8
comps  9 comps 10 comps
# CV      1.5e+07 3418008 2414639 2049701 1806343 1371151 1325548 1301942
1297801 1283053 1282432
# adjCV    1.5e+07 3417887 2414598 2049563 1805870 1370529 1325245 1301788
1297869 1282936 1282273
#      11 comps 12 comps 13 comps 14 comps 15 comps 16 comps 17 comps
# CV      1280546 1279732 1278057 1277875 1277232 1276966 1276987
# adjCV    1280398 1279576 1277894 1277712 1277070 1276804 1276825

pls.pred <- predict(pls.fit, train.data, ncomp=which.min(MSEP(pls.fit)$val[1,1,]) - 1)
mean((Apps - pls.pred)^2) # training error
# [1] 1059279

```